

## Homework 1: Crawling

Released: Jan 16<sup>th</sup>, 2020Due: Jan 23<sup>rd</sup>, 2020 @ 23:59

### Ground Rules

This homework must be done individually. You can ask others for help with the tools, however, the submitted homework has to be your own work.

### Summary

In this homework, you will utilize web crawlers to collect webpages and extract data from the Internet Movie Database (IMDb) website (<https://www.imdb.com/>).

A web crawler is a program/bot that systematically browses the World Wide Web (WWW), typically for the purpose of web indexing (web spidering). It starts with a list of seed URLs to visit, and as it visits each webpage, it finds the links in that web page, and then visits those links and repeats the entire process. There are two web crawler technologies you are required to use in this homework:

- Scrapy (<https://scrapy.org/>): A python library crawler
- import.io (<https://import.io/>): An interactive crawling platform

### Task 1: Scrapy (7 points)

#### Task 1.1 (3 points)

Crawl at least **5000** webpages of **Sci-Fi movies/shows** in IMDb using Scrapy. Extract and generate the following attributes for each webpage:

id	unique id for the webpage
url	url of the webpage
timestamp_crawl	timestamp of the crawling event
title	see Figure 1 if attribute doesn't exist, set as empty string
release_date	
budget	
gross_usa	
runtime	

Store your crawled data into a JSON-Lines (.jl) file. In this file format, each line is a valid JSON object (dictionary) that holds the attributes listed above for a single crawled webpage. You can check the attached file `sample.jl` to understand the format. While crawling, please make sure you obey the website's politeness rules (i.e. sleep time between requests) in order to avoid getting banned.

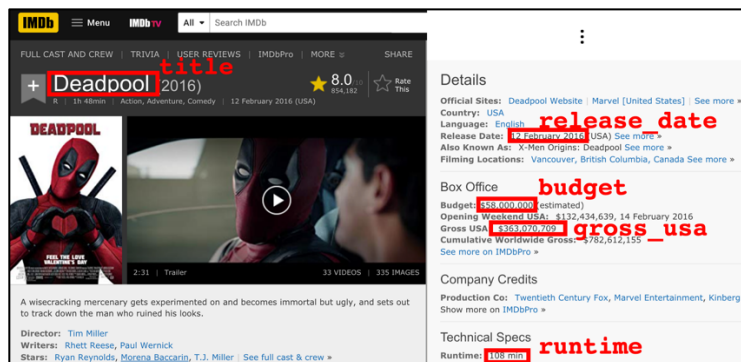


Figure 1: An example movie/show webpage with the required attributes

### Task 1.2 (3 points)

Similar to the previous task, crawl at least **5000** webpages of cast (actors and actresses) in IMDb using Scrapy. Extract and generate the following attributes for each cast webpage:

id	unique id for the webpage
url	url of the webpage
timestamp_crawl	timestamp of the crawling event
name	
date_of_birth	see Figure 2
place_of_birth	
date_of_death	if attribute doesn't exist, set as empty string
place_of_death	

Similarly, store your crawled data into a JSON-Lines (`.jsonl`) file.

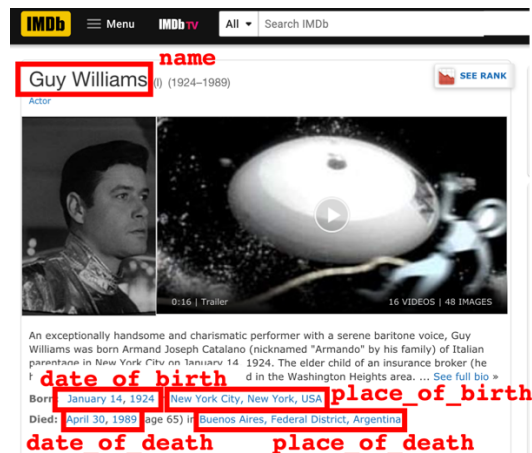


Figure 2: An example cast webpage with the required attributes

### Task 1.3 (1 points)

Answer the following questions (no more than 2 sentences for each question):

- 1.3.1 What is the seed URL(s) you used for each task?
- 1.3.2 How did you manage to only collect movie/show or cast pages?
- 1.3.3 Did you need to discard irrelevant pages? If so, how?
- 1.3.4 Did you collect the required number of pages? If you were not able to do so, please describe and explain your issues.

### Task 2: import.io (3 points)

Use the interactive platform of **import.io** to crawl and extract the same attributes as in Task 1.1 from a total of **400 Romance movies/shows** in IMDb. Store your output in a comma-separated values (`.csv`) file. You are expected to create a chain of at least 2 chained extractors, in which at least 1 extractor's output is used to generate the seed URLs of another extractor (e.g. the movie/show webpage extractor). The general expected chain pipeline is shown in Figure 3. Describe the chain you constructed. For each extractor in the chain: describe its inputs, outputs and role (what is it doing) and attach a screenshot of the run including the top output rows (by clicking "Preview Data" in the "Run History" screen), as seen in Figure 4.

- The free trial version of the service is limited to 1000 URL queries, which should be sufficient for the chain of extractors you are expected to construct for this task.
- You may use the "URL Generator" to create the list of seed URL(s) of your first extractor in the chain. This is not mandatory but may make the task easier for you.




Figure 3: The general extraction pipeline you are expected to build



Figure 4: A screenshot example of an extractor's run

## Submission Instructions

You must submit (via Blackboard) the following files/folders in a single `.zip` archive named `Firstname_Lastname_hw01.zip`:

- `Firstname_Lastname_hw01_report.pdf`: pdf file with your answers to Task 1.3 and 2
- JSON-Lines files containing the data you crawled using Scrapy for Task 1.1 and 1.2:
  - `Firstname_Lastname_hw01_scrapy_title.jl`: Generated data from Task 1.1
  - `Firstname_Lastname_hw01_scrapy_cast.jl`: Generated data from Task 1.2
- `Firstname_Lastname_hw01_importio_title.csv`: A Comma-Separated values file of the generated data results from Task 2
- `source`: This folder includes all the code you wrote to accomplish Task 1 (i.e. your Scrapy crawler, seed files, your script/program to eliminate unwanted pages and store webpages into JSON-Lines format, etc...)

## Additional Notes

1. We also provide you with a script called `post_processing.py`, you can use this script to validate the structure of your JSON-Lines files in Task 1.1:
  - a. Prerequisites: python 3 and the `ujson` package (`pip install <package>`)
  - b. Usage: The script takes one argument `<path>` which is the path of your `jl` file, processes the file and prints a message to let the user know if it is valid.  
`python post_processing.py /home/users/sample.jl` will print:  
 Process: (0/2)...  
 Process: (1/2)...  
 Finished processing, looks good, found 2 entries.
2. It is your responsibility to validate the structure of the outputs in tasks 1.2 and 2 prior to submitting.